

Aarhus University

Decision Support Systems

PROJECT REPORT

Author:

Supervisor:

Lasse Lildholdt
(201507170)
Stinus Skovgaard
(201507170)
Daniel Tøttrup
(201507170)
Johan Vasegaard
(201507170)
Frederik Madsen
(201504477)

Christian Fischer Pedersen

11. maj 2020

Indhold

1	Introduction	1
2	Simple linear regression / Multiple linear regression	2
2.1	Predict β_0 and β_1	2
2.2	Multiple linear regression	5
2.2.1	Estimating the regression constants	5
2.3	Lab results	7
2.3.1	Lab 3.6.2	7
2.3.2	Lab 3.6.3	8
3	Logisitic regression / Linear discriminant analysis	9
4	Cross validation / Bootstrap	10
5	Subset selection	11
5.1	Best subset selection	11
5.2	Stepwise Selection	12
5.2.1	Forward stepwise selection	12
5.2.2	Backward stepwise selection	13
5.3	Choosing the optimal model	13
5.4	Lab 6.5.1	15
6	Shrinkage And Dimension Reduction Methods	17
6.1	Ridge regression	17
6.2	The Lasso	18
6.3	Principal Component Regression	18
6.4	Partial Least Squares	19
6.5	Lab results	19
7	Clustering Methods	22
7.1	K-means clustering	22
7.1.1	Hierarchical Clustering	24
7.2	Lab results	24

Figurer

2.1	Advertisement data	3
2.2	Linear regression on TV advertisement	3
2.3	The regression pane with TV and Radio as predictor variables	6
2.4	Plot of data and regression line	7
2.5	Properties of regression line	8
5.1	Lab 6.5.1 plotted results	16
6.1	PCA on the dataset "Hitters"	20
6.2	MSE as a function of components in PCR	20
6.3	MSE as a function of components in PLS	21
7.1	K-means clustering example	22
7.2	K-means clustering pseudo code	23
7.3	K-means clustering in action	23
7.4	Dendrogram example	24
7.5	Complete histogram of clusters in the data set	25
7.6	Clusters generated by KMeans with $k = 2$ and $k = 3$	25

INTRODUCTION

In this project report, the reader will be presented with problem solutions for the course Decision support systems. Throughout the solution the reader will achieve knowledge on several different subject within the main area. Each topic will be presented with the theory along with solutions for appropriate exercise to validate the presented theory.

The main topics which will be handled in this report will be as follows:

- Simple linear regression / Multiple linear regression
- Logisitic regression / Linear discriminant analysis
- Cross validation / Bootstrap
- Subset selection
- Shrinkage methods / DIMension reduction methods

SIMPLE LINEAR REGRESSION / MULTIPLE LINEAR REGRESSION

The simple Linear Regression approach is a quick and simple method for predicting a response Y based on X . The linear model is used to give an idea of the relationship between to dataset. For this to be true an assumption that the two variables have a linear relationship is needed. The mathematical representation of this can be seen below.

$$Y \approx \beta_0 + \beta_1 X \quad (2.1)$$

eq. (2.1) can also be seen as "Regressing Y onto X ". As an example the dataset Advertising.csv contains sales on a certain product and advertisement money spent on certain media platforms. X represents TV advertising and Y represents sales. It is the possible to regress sales onto TV. This is expressed as:

$$sales \approx \beta_0 + \beta_1 TV \quad (2.2)$$

The two constants β_0 and β_1 represents the intercept (Where it intercepts the y-axis) and slope of the linear model. These constant needs to be predicted and when they have we will end out with a linear model that fits our data.

2.1 Predict β_0 and β_1

In the Advertising.csv dataset a number of observations ($n = 200$) have been made on amount of advertisements for tv, radio, newspaper and the corresponding sales. The advertisement data can be seen on fig. 2.1.




Figure 2.1: Advertisement data

To estimate β_0 and β_1 we need to introduce RSS which is the *residual sum of squares*. RSS is the difference or error between observed response value and predicted response value that is predicted by our linear model. On fig. 2.2 a linear model is put on top of the TV advertisement data. The RSS is the summed value of the difference between the red points to the blue line. We want this value to be as low as possible to get the most accurate model.




Figure 2.2: Linear regression on TV advertisement

We can define the RSS value as:

$$\begin{aligned}
 RSS &= e_1^2 + e_2^2 \dots e_n^2 \\
 &\text{where} \\
 e &= y_i - \hat{y}
 \end{aligned}
 \tag{2.3}$$

This means that to get to a final method we need to minimize RSS. Some calculus shows that these minimizers are:

$$\begin{aligned}
 \hat{\beta}_1 &= \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} \\
 \beta_0 &= \bar{y} - \beta_1 \bar{x} \\
 &\text{where} \\
 \bar{y} &= \frac{1}{n} \sum_{i=1}^n y_i \\
 &\text{and} \\
 \bar{x} &= \frac{1}{n} \sum_{i=1}^n x_i
 \end{aligned}
 \tag{2.4}$$

This method of estimating the constants have been used on fig. 2.2. How this is done can be seen on the code snippet below.

```

x_mean = get_mean(x_data)
y_mean = get_mean(y_data)

num = 0
den = 0

for i in range(len(x_data)):
    num += (x_data[i]-x_mean)*(y_data[i]-y_mean)
    den += (x_data[i] - x_mean) ** 2

b1_hat = num/den
b0_hat = y_mean-b1_hat*x_mean

return b0_hat, b1_hat

```

β_0 and β_1 is estimated to 7.03 and 0.04 respectively for the TV advertisement data.

There will always be an error of some sort when estimating this linear model. Because of this we can add a error constant to eq. (2.1). The error constant (epsilon) is generated based on a gaussian distribution with a mean of 0. The equation will look as the following:

$$Y \approx \beta_0 + \beta_1 X + \epsilon \tag{2.5}$$

2.2 Multiple linear regression

Simple linear regression is a great to see the response of a single predictor variable. However there are times where it would be beneficial to have multiple predictor variables. Instead of having only a 1-dimensional regression line we can with multiple linear regression have a many-dimensional regression line (plane for 2D, cloud for 3D).

In simple linear regression we had an intercept constant, β_0 and a slope constant β_1 . In this model we essentially have the same plus some extra slope constant and predictor variables. This can be seen on eq. (2.6).

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \epsilon \quad (2.6)$$

If we fit this equation to our advertisement example we would get:

$$\text{sales} = \beta_0 + \beta_1 \mathbf{TV} + \beta_2 \mathbf{Radio} + \beta_3 \mathbf{Newspaper} + \epsilon \quad (2.7)$$

2.2.1 Estimating the regression constants

To estimate the constants we do same as in the simple regression method with some changes. We still try to minimize the RSS value, but since there are more constants the RSS value is calculated with the equation eq. (2.8). However to estimate the constants itself the use of a python package has been used. Sklearn has a regression module that can estimate our constants. This can be seen in the codes snippet further down in this chapter.

$$\begin{aligned} RSS &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_{i1} - \hat{\beta}_2 x_{i2} - \dots - \hat{\beta}_p x_{ip})^2 \end{aligned} \quad (2.8)$$

When using two predictor variables the regression line becomes a plane. On fig. 2.3 the regression plane when using TV and Radio as predictor variables is shown. The plane is placed in a way such that the RSS value is as low as possible.




Figure 2.3: The regression pane with TV and Radio as predictor variables

The code for making the regression plane is listed below:

```
# Multiple Regression
regr = skl_lm.LinearRegression()

X = advertising[['Radio', 'TV']].values
y = advertising.Sales

# Fitting data to estimate constants.
regr.fit(X, y)

# Make it so that the plane fits all points
Radio = np.arange(0, 50)
TV = np.arange(0, 300)

B1, B2 = np.meshgrid(Radio, TV, indexing='xy')
Z = np.zeros((TV.size, Radio.size))

for (i, j), v in np.ndenumerate(Z):
    # The response on TV and Radio
    Z[i, j] = (regr.intercept_ + B1[i, j] * regr.coef_[0] + B2[i, j] * regr.coef_[1])
```

2.3 Lab results

2.3.1 Lab 3.6.2

The data used in this lab exercise is from the MASS library which contains housing data in the Boston area. First a linear regression model will be applied on the "medv" and "lstat". **medv** as response (y) and **lstat** as the predictor. The code snippet below shows how β_0 and β_1 is estimated. This linear model is plotted on the fig. 2.4.

```
x_mean = get_mean(x_data)
y_mean = get_mean(y_data)

num = 0
den = 0

for i in range(len(x_data)):
    num += (x_data[i]-x_mean)*(y_data[i]-y_mean)
    den += (x_data[i] - x_mean) ** 2

b1_hat = num/den
b0_hat = y_mean-b1_hat*x_mean

return b0_hat, b1_hat
```




Figure 2.4: Plot of data and regression line

To get a more detailed look at the different statistical properties, the following code is executed and the use of sklearn's LinearRegression module is used.

```

regr = skl_lm.LinearRegression()
regr.fit(X,y)

params = np.append(regr.intercept_, regr.coef_)
predictions = regr.predict(X)

newX = pd.DataFrame({"Constant": np.ones(len(X))}).join(pd.DataFrame(X))
MSE = (sum((y - predictions) ** 2)) / (len(newX) - len(newX.columns))

var_b = MSE * (np.linalg.inv(np.dot(newX.T, newX)).diagonal())
sd_b = np.sqrt(var_b)
ts_b = params / sd_b

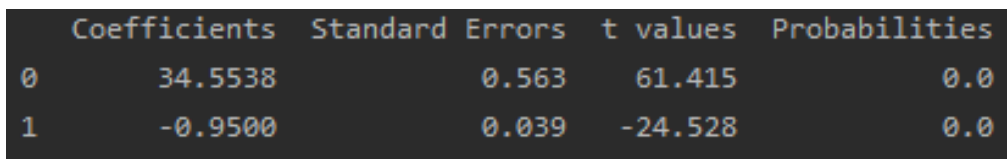
p_values = [2 * (1 - stats.t.cdf(np.abs(i), (len(newX) - 1)))] for i in ts_b]

sd_b = np.round(sd_b, 3)
ts_b = np.round(ts_b, 3)
p_values = np.round(p_values, 3)
params = np.round(params, 4)

myDF3 = pd.DataFrame()
myDF3["Coefficients"], myDF3["Standard Errors"], myDF3["t values"], myDF3["Probabilities"] = [params, sd_b, ts_b,
p_values]
print(myDF3)

```

This code will output the following:



	Coefficients	Standard Errors	t values	Probabilities
0	34.5538	0.563	61.415	0.0
1	-0.9500	0.039	-24.528	0.0

Figur 2.5: Properties of regression line

2.3.2 Lab 3.6.3

Now we want to fit multiple predictor variables using the multiple linear regression method. This is also done using Sklearn LinearRegression module. The statistical properties when using all 13 predictor variables can be seen on ??

	Coefficients	Standard Errors	t-values	Probabilities
0	36.4595	5.103	7.144	0.000
1	-0.1080	0.033	-3.287	0.001
2	0.0464	0.014	3.382	0.001
3	0.0206	0.061	0.334	0.738
4	2.6867	0.862	3.118	0.002
5	-17.7666	3.820	-4.651	0.000
6	3.8099	0.418	9.116	0.000
7	0.0007	0.013	0.052	0.958
8	-1.4756	0.199	-7.398	0.000
9	0.3060	0.066	4.613	0.000
10	-0.0123	0.004	-3.280	0.001
11	-0.9527	0.131	-7.283	0.000
12	0.0093	0.003	3.467	0.001
13	-0.5248	0.051	-10.347	0.000

Figur 2.6: Properties of multiple regression line

LOGISITIC REGRESSION / LINEAR DISCRIMINANT ANALYSIS

CROSS VALIDATION / BOOTSTRAP

SUBSET SELECTION

This chapter will address the subject Subset selection. Subset selection concerns with selecting or shrinking the coefficients of features to make the model more interpretable and in some cases to predict better.

Linear models are simple and can be interpreted because it usually has a small number of coefficients. In cases where the number of predictors is bigger than the number of samples, we can't use the full least squares, because the solutions is not even defined. In such cases we must reduce the number of features to be able to obtain a solution. It is also very important to not fit your data too hard which regularize, or selection of features also helps with. Along the same lines, when we have a small number of features the model becomes more interpretable.

There exist many different methods to perform subset selection. One of the methods to select the most important features regarding a specific response is called best subset selection. This is where we identify a subset of the predictors that is most related to the response.

5.1 Best subset selection

Best subset selection algorithm is a simple algorithm used to understand which predictors are mostly linked to the responds. To perform best subset selection, we fit a separate least squares regression for each possible combination. It starts out by fitting all p models that contain exactly one predictor, then fitting all p models that contain exactly two predictors and so forth. The number of models to fit can be calculated like this:

$$\binom{p}{k} = \frac{p!}{k!(p-k)!} \quad (5.1)$$

Where p is the number of all predictors, and k is the number of predictors in the subset. After all possible combinations of p models has been fitted, we then look at the resulting model, with the goal of identifying the best one.

Best subset selection algorithm can be divided into three steps:

1. Let M_0 denote the *null model*, which contain no predictors. This model simply predicts the sample mean for each observation.
2. For $k=1,2,\dots,p$:
 - (a) Fit all $\binom{p}{k}$ models that contain exactly k predictors.

- (b) Pick the best among these $\binom{p}{k}$ models, and call it M_k . Here *best* is defined as having the smallest RSS, or equivalently largest R^2
3. Select a single best model from among M_0, \dots, M_p using cross-validated prediction error, AIC, BIC or adjusted R^2

The task of selecting the best subset model, must be performed with care, because RSS decreases monotonically and the R^2 increases monotonically, as the number of features included in the models increases. So, if we use these statistics, we will always end up with the model involving all the variables. Another problem with a low RSS or a high R^2 is that it only indicates a model with a low training error. And a low training error does not equal a low test error. So, for those reasons we need to use cross-validation, AIC, BIC or adjusted R^2 .

5.2 Stepwise Selection

If the number of predictors p gets too large, best subset selection algorithm cannot be applied for computational reasons. Another reason why best subset selection is not always the best algorithm to select a subset, is for statistical reasons if the p is too large. Then there is a higher chance of finding models that look good on the training data, even though they might not have any predictive power on future data. For these reasons, stepwise methods, which explore a far more restricted set of models, are attractive alternatives to best subset selection.

In this section we'll touch upon two different stepwise selection methods; the forward stepwise selection method and the backwards stepwise selection method.

5.2.1 Forward stepwise selection

Forward stepwise selection method is very similar to the best subset selection method. Like the best subset selection method, the forward stepwise selection also begins with a model containing no predictors, and adds one predictor at a time until all predictors are in the model. But unlike the best subset selection method the forward stepwise selection doesn't look at all possible models that contain k predictors at each step. Instead, we are just looking at the models that contain the $k-1$ predictors that already were chosen in the previous step, plus one more. This means that at the k -th step, we are looking at a much more restricted set of models compared to the best subset selection method.

Forward stepwise selection method can be divided into three steps:

1. Let M_0 denote the *null model*, which contain no predictors.
2. For $k=0, \dots, p-1$:
 - (a) Consider all $p-k$ models that augment the predictors in M_k with one additional predictor.

- (b) Chose the best among these $p-k$ models, and call it $M_k + 1$. Here *best* is defined as having the smallest RSS, or highest R^2
3. Select a single best model from among M_0, \dots, M_p using crass-validated prediction error, AIC, BIC or adjusted R^2

Compared to best subset selection forward stepwise selection has a computational advantage, as we consider 2^p models in best subset selection and only p^2 models in forward stepwise selection method. This relates to the fact that forward stepwise selection is not guaranteed to find the best possible model out of all 2^p models containing subset of p predictors.

5.2.2 Backward stepwise selection

Backward stepwise selection is exactly the opposite of forward stepwise selection. In contrast backward stepwise selection begins with the full least squares model containing all p predictors, and then removes the least useful predictor, one at a time.

Forward stepwise selection method can be divided into three steps:

1. Let M_p denote the *full model*, which contain all p predictors.
2. For $k=p, p-1, \dots, 1$:
 - (a) Consider all k models that contain all but one of the predictors in M_k , for a total of $k-1$ predictors.
 - (b) Chose the best among these k models, and call it $M_k - 1$. Here *best* is defined as having the smallest RSS, or highest R^2
3. Select a single best model from among M_0, \dots, M_p using crass-validated prediction error, AIC, BIC or adjusted R^2

Just like forward stepwise selection, backwards stepwise selection is not guaranteed to gives us the best model containing a particular subset of p predictors.

The major difference between forward and backwards stepwise selection is that the number of samples needs to be larger than the number of variables to perform the backward stepwise selection method (so its possible to fir a least squares model). This is not case for forward stepwise selection, it can be applied when $n < p$ and $p > n$.

5.3 Choosing the optimal model

As already explained RSS and R^2 are not suitable for selecting the best model among a collection of models, because these quantities are related to the training error and not the test error. So, in order to select the best model with respect to the test-error there are two common approaches:

1. Indirectly estimate the test error by making an adjustment to the training error to account for the bias due to overfitting.
2. Directly estimate the test error by either using a validation set approach or a cross-validation approach.

In this section we'll only talk about AIC, BIC and adjusted R^2 which all are indirectly estimate of the test error.

AIC

AIC stand for Akaike information criterion and deals with the trade-off between goodness of fit of the model and simplicity of the model or it deals with the risk of overfitting and underfitting.

$$AIC = \frac{1}{n\hat{\sigma}^2}(RSS + 2d\hat{\sigma}^2) \quad (5.2)$$

The best model according to the AIC is where the AIC is smallest.

BIC

BIC stands for Bayesian information criterion. This is closely related to the AIC. Both the BIC and the AIC introduces a penalty term for number of parameters in the model to resolve the problem of overfitting. The penalty term is larger in the BIC than in the AIC.

$$BIC = \frac{1}{n}(RSS + \log(n)d\hat{\sigma}^2) \quad (5.3)$$

The best model according to the BIC is where the BIC is smallest.

Adjusted R^2

Another very popular approach for selecting among a set of models that contain a different number of variables. R^2 is defined as $1 - RSS/TSS$, where TSS is the total sum of squares for the response. But as already mentioned R^2 keeps increasing as more variables are added to the model. The adjusted R^2 is calculated as:

$$AdjustedR^2 = 1 - \frac{RSS/(n - d - 1)}{TSS/(n - 1)} \quad (5.4)$$

Unlike the AIC and BIC, a large value for adjusted R^2 indicates a model with a small test error.

5.4 Lab 6.5.1

In Lab 6.5.1 we have implemented best subset selection and performed it on the “Hitters”-dataset, to predict a baseball players salary based on various statistics. This has been implemented in python and will be discussed briefly in this section, for a more detailed look into the implementation see the source-code in the appendix.

In this section only the heart of the implemented algorithm will be presented. The function `getBestModel` returns the model with the lowest RSS, based on all combinations of k number of features.

```
def getBestModel(k, X, y):
    results = []
    # Goes through all combinations of k number of features
    for combination in itertools.combinations(X.columns, k):
        results.append(subset_Process(combination, X, y))

    # Stores all combinations in a single dataframe
    models = pd.DataFrame(results)

    # Select the model with the lowest RSS
    best_model = models.loc[models['RSS'].argmin()]

    return best_model
```

The function `subset_Process` is called by the `getBestModel` function described above. This function performs linear regression on a model and calculates the RSS the given model.

```
def subset_Process(predictor, X, y):
    # sm.OLS is an estimator that performs linear regression on a model
    temp_model = sm.OLS(y, X[list(predictor)])
    model = temp_model.fit()

    # Then calculate the RSS for the chosen model, and return the model and its RSS together
    RSS = ((model.predict(X[list(predictor)]) - y) ** 2).sum()
    return {"model": model, "RSS": RSS}
```

These two functions are repeated a lot of times to estimate all possible combinations of k predictors.

On fig. 5.1 are the results plotted. Because the number of combinations increases dramatically as k increases so does the computation time. For this reason, k is equal to nine. We can see that according to the BIC, the model with 6 variables performs the best. But according to Adjusted R^2 and AIC a model with more variables than six might be better. Again none of these measures gives us an entirely accurate picture, but they all agree that a model with fewer than five predictors is insufficient.

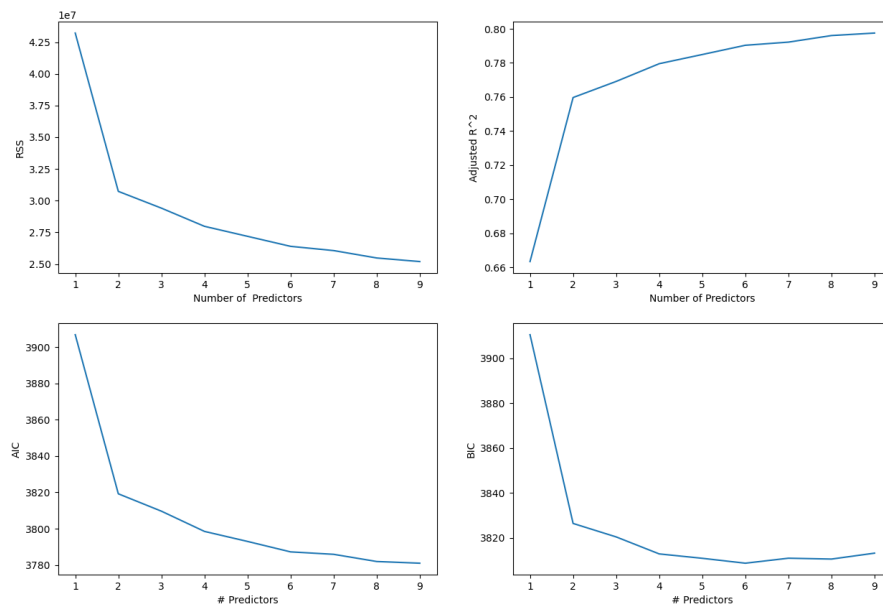


Figure 5.1: Lab 6.5.1 plotted results

SHRINKAGE AND DIMENSION REDUCTION METHODS

Contrary to subset selection methods, which uses least squares to fit a linear model containing a subset with n of all p predictors, shrinkage methods fits a model containing all p predictors. This is done by a penalty that regularizes the coefficient estimates and thereby shrinks them towards zero. By shrinking the coefficient estimates their variance can be significantly reduced. In this section the two shrinkage methods Ridge regression and The Lasso will be covered.

Both shrinking methods tries to control variance by either using a subset of the original variables or by shrinking their coefficients towards zero and uses all of the original predictors. Another class of approaches, dimension reduction, is one that transforms the predictors first and then fits a least squares model using the transformed predictors. The reduction comes from the fact that the methods reduces the problem of estimating $p + 1$ coefficients to estimating $M + 1$ where $M < p$. In this section the two dimension reduction methods Principal Component Regression and Partial Least Squares will be covered.

6.1 Ridge regression

Ridge regression is very similar to least square fitting in that it seeks to minimize RSS, but it adds a second term called the shrinkage penalty.

The equation 6.1 shows the full equation for ridge regression. Here the first term is the RSS where β_0, \dots, β_p is to be estimated such that it is minimized, but the second term introduces a penalty to β_j which effectively shrinks it towards zero. This penalty is scaled with λ such that as it moves towards zero, the penalty moves towards zero and the equation produces the least squares estimates. Moving the tuning parameter towards infinity will in turn drive the coefficient estimates towards zero.

Thus ridge regression will produce a different set of coefficient estimates, $\hat{\beta}_\lambda^R$, for each value of λ hence making it a tuning parameter. It is then critical to choose a good value for λ , which can be done e.g. by the cross-validation method.

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2 \quad (6.1)$$

The advantage of ridge regression over standard least squares comes from the bias-variance trade-off. By increasing the value of λ and thereby the effect of the penalty term, it is possible to decrease the variance of the predictor however the bias increases. As the test MSE (mean squared error) is a function of the variance plus the squared

bias, finding a λ value which decreases the variance more than it increases the bias can result in a lower test MSE. Thus ridge regression will be superior whenever the least squares estimates have high variance.

6.2 The Lasso

The Lasso improves upon a disadvantage of ridge regression, namely that it includes all p predictors in the final model, because it only reduces the magnitudes of the coefficients but never actually excludes any of them. Where ridge regression uses the ℓ_2 -norm in its regularization term (equation 6.1), the Lasso uses ℓ_1 -norm (equation 6.2). This has the effect of forcing some of the coefficients to be zero, depending on the value of λ , instead of only driving them towards zero. Therefore the Lasso acts like subset selection, as it effectively performs variable selection yielding a sparse model which exactly is its improvement over ridge regression.

Selecting a good value for λ is critical just as it is for ridge regression and can be done by cross-validation.

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p |\beta_j| \quad (6.2)$$

6.3 Principal Component Regression

As with all dimension reduction methods Principal Component Regression (PCR) works in two steps: first transformed predictors are obtained and then a model is fit using the transformed predictors.

The underlying method for obtaining the transformed predictors is in this case PCA (Principal Component Analysis). PCA derives a low-dimensional set of features from a large set of variables by acquiring basis vectors, *principal components*, that forms an orthogonal basis. This is done by finding vectors with values that minimize the sum of squared perpendicular distances between each point and the vector. Finding the next principal component is done as a linear combination of the variables that is uncorrelated with the principal component before it and has the largest variance. Up to $M \leq p$ principal components can be constructed this way, with p being the number of predictors. By this construction the first principal component will contain the most information of the data-set with each following principal component containing less and less information. Fitting a least square model to Z_1, \dots, Z_M principal components instead of X_1, \dots, X_p data points with $p \ll M$, can in theory lead to better results as it can mitigate overfitting. This stems from the notion that most or all of the information in the data relating to the response is contained in the Z_1, \dots, Z_M principal component. If however $p = M$ PCR will perform the same as doing least square fitting on all of the original predictors. In PCR the number of principal components M used for least square fitting becomes a hyper parameter that needs to be chosen carefully. This is typically done using the cross-validation method.

6.4 Partial Least Squares

In PCR the M principal components is guaranteed to best explain the predictors, however they are not guaranteed to best explain the response. This is due to the unsupervised nature of PCR, where the response does not supervise the identification of the transformed features (principal components).

Unlike PCR, Partial Least Squares (PLS) is supervised in the sense that the response is used to identify the transformed features that not only approximates the original predictors but also relates to the response.

PLS first identifies a new set of features Z_1, \dots, Z_M that represents $M < p$ linear combinations of the original p predictors. This is done by first standardizing the p predictors. Then the first translated feature Z_1 is calculated as in equation 6.3, with the difference that each θ_{j1} is set equal to the coefficient from the simple linear regression of the response Y onto X_j . In doing this tweaked version PLS places a higher weight on variables that are strongly related to the response. To compute Z_2 the residuals from regressing each variable on Z_1 are then used following the same calculations as for determining Z_1 . Lastly least squares as in equation 7.1 is used to fit a linear model to predict Y using Z_1, \dots, Z_M translated feature instead of the original p features just like in PCR.

$$Z_m = \sum_{j=1}^p \phi_{jm} X_j \quad (6.3)$$

$$y_i = \theta_0 + \sum_{m=1}^M \theta_m Z_{im} + \epsilon_i \quad (6.4)$$

6.5 Lab results

Principal Component Regression(PCR) and Partial Least Squares(PLS) has been applied on the data set "Hitters" to predict salary as in exercise 6.7.1 and 6.7.2 in the course book. This has been done by finding the number of components that gives the smallest MSE on the training set.

Tenfold cross validation has been used to find the MSE as a function of components as seen on code listing 6.5 showing the implementation for PCR. A similar implementation has been used for PLS. The cross validation has been done for 10 components, as PCA shows an explained variance of over 90% with a flattened curve beyond 10 components, as seen on figure 6.1.

The results for PCR can be seen on figure 6.2, where six principal components gives the lowest MSE. For PLS the lowest MSE value is achieved using only two components. Applying PCR with seven components on the test set gives a MSE of 114098.06, while applying PLS with two components on the test set gives a MSE of 104838.51. Thus for the data set used in this example PLS clearly performs better with both a lower component count and a lower MSE value on the test set.

```
def subset_Process(predictor, X, y):
```

```

# Scale the data
pca = PCA()
X_reduced = pca.fit_transform(scale(X_train))
# Setup folds and linear regression
mse = []
regr = linear_model.LinearRegression()
kf_10 = KFold(n_splits=10, shuffle=True, random_state=1)
# Single CV to get MSE for the intercept
score = cross_val_score(regr, np.ones((len(X_reduced),1)), y.ravel(), cv=kf_10, scoring='neg_mean_squared_error').mean()
mse.append(-score)
# CV for pc principal components
for i in np.arange(1, pc):
    score = cross_val_score(regr, X_reduced[:,i], y, cv=kf_10, scoring='neg_mean_squared_error').mean()
    mse.append(-score)
# Plot MSE as function of principal components after linear regression
....

```

Code Listing 6.1: Code snippet showing tenfold cross validation of PCR




Figure 6.1: PCA on the dataset "Hitters"




Figure 6.2: MSE as a function of components in PCR




Figure 6.3: MSE as a function of components in PLS

CLUSTERING METHODS

Clustering refers to grouping data into different clusters or classes. We do this to get observations or samples clustered into groups where they share some similarities. This is used extensively in the field of machine learning. An example of clusters or classes could be handwritten numbers (0-9). Each number has distinct features which can be divided into clusters. An unsupervised problem would be to divide the numbers into subclasses. In a supervised problem we would already have made the clustering and then divide new samples into the already known clusters, or in other words, predict some outcome based on the sample.

7.1 K-means clustering

K-means clustering is a simple method of clustering data. It works by having a number (K) mean-vectors and assigning a sample to the mean-vector that it's closest to. On fig. 7.1 150 sample have been generated and three different values for K. It is easy to see that the K-means method tries to cluster samples that are close to each other in the same class.




Figure 7.1: K-means clustering example

This works by minimizing the squared euclidean distance to each mean-vector.

$$\text{minimize}_{C_1 \dots C_K} \left\{ \sum_{k=1}^K \frac{1}{|C_k|} \sum_{i, i' \in C_k} \sum_{j=1}^p (x_{ij} - x_{i'j})^2 \right\} \quad (7.1)$$

$|C_k|$ denotes the number of samples in the k th cluster. Or the within variation of the k th cluster which is the sum of all the pairwise squared euclidean distances between the samples in the k th cluster, divided by the total number of samples in k th cluster.


To implement this as an algorithm the book [?] covers how this should be done.

Algorithm 10.1 *K-Means Clustering*

1. Randomly assign a number, from 1 to K , to each of the observations. These serve as initial cluster assignments for the observations.
 2. Iterate until the cluster assignments stop changing:
 - (a) For each of the K clusters, compute the cluster *centroid*. The k th cluster centroid is the vector of the p feature means for the observations in the k th cluster.
 - (b) Assign each observation to the cluster whose centroid is closest (where *closest* is defined using Euclidean distance).
-

Figur 7.2: K-means clustering pseudo code

In action this algorithm performs well and can be seen on ??



Figur 7.3: K-means clustering in action

7.1.1 Hierarchical Clustering

Hierarchical Clustering makes use of tree-based representation of the samples. This tree-based representation is called a Dendrogram and looks like an upside-down tree. An example of a dendrogram can be seen on fig. 7.4 [?]




Figure 7.4: Dendrogram example

Its easy to see how the clustering works. It starts from the bottom works its way up. To see if two samples are similar we can compare how far up on the vertical axis two samples fuses. If they fuse at the bottom they are very similar and at to top not very similar.

A hierarchical clustering algorithm will work like the dendrogram. That means that if we have n sample, it will also starts at the bottom and look at the two samples that are most similar and fuses them together so we have $n-1$ clusters. Next it will again fuse the two most similar samples and will be left with $n-2$ clusters. The algorithm will keep going until all samples belong to a cluster.

7.2 Lab results

The data for the two lab exercises 10.5.1 and 10.5.2 are generated with the code listed in 7.2. The data is made so that it contains two natural clusters.

Performing KMeans on the generated data with $k = 2$ and $k = 3$ both with 20 random initial assignments results in the data clustering on figure 7.6. For $k = 2$ the data is clearly divided into two clusters (red and green), while for $k = 3$ it is divided into three clusters (red, green and yellow). Adding additional clusters also results in different cluster means as is evident by the black +, which denotes the given clusters mean.

Figure 7.5 shows the "complete" hierarchical clustering of the data, which is generated as "bottom up" clustering by comparing the largest pairwise distance between points in possible clusters. The colors red and green marks the two clusters that is known to be in the data, while the blue color marks the "single cluster" being all the data.

Using "single" hierarchical clustering of the data, the graph would instead show clustering by comparing the smallest pairwise distance between points in possible clusters.

Lastly also "average" hierarchical clustering could be used, which of course would use the average distance instead.

```
# Generate data
np.random.seed(2)
X = np.random.standard_normal((50,2))
X[:25,0] = X[:25,0]+3
X[:25,1] = X[:25,1]-4
```

Code Listing 7.1: Code for generating the data




Figure 7.5: Complete histogram of clusters in the data set




Figure 7.6: Clusters generated by KMeans with $k = 2$ and $k = 3$